

```
Activities Terminal Mar 23 16:15 samara@hitloop: ~/Videos

samara@hitloop: ~/Videos samara@hitloop: ~/Videos samara@hitloop: ~/Videos samara@hitloop: ~/Videos

(base) samara@hitloop:~/Videos$ uvicorn app:app --host 127.0.0.1 --port 8000 --reload
INFO: Will watch for changes in these directories: ['/home/samara/Videos']
INFO: Uvicorn running on http://127.0.0.1:8000 (Press CTRL+C to quit)
INFO: Started reloader process [730034] using WatchFiles
2025-03-23 13:46:41.549887: I tensorflow/core/util/port.cc:153] oneDNN custom operations are on. You may see slightly different numerical results due to floating-point round-off errors from different computation orders. To turn them off, set the environment variable `TF_ENABLE_ONEDNN_OPTS=0`.
2025-03-23 13:46:41.558022: E external/local_xla/xla/stream_executor/cuda/cuda_fft.cc:467] Unable to register cuFFT factory: Attempting to register factory for plugin cuFFT when one has already been registered
WARNING: All log messages before absl::InitializeLog() is called are written to STDERR
E0000 00:00:1742717801.567849 730036 cuda_dnn.cc:8579] Unable to register cuDNN factory: Attempting to register factory for plugin cuDNN when one has already been registered
E0000 00:00:1742717801.571078 730036 cuda_blas.cc:1407] Unable to register cuBLAS factory: Attempting to register factory for plugin cuBLAS when one has already been registered
W0000 00:00:1742717801.579269 730036 computation_placer.cc:177] computation placer already registered. Please check linkage and avoid linking the same target more than once.
W0000 00:00:1742717801.579281 730036 computation_placer.cc:177] computation placer already registered. Please check linkage and avoid linking the same target more than once.
W0000 00:00:1742717801.579283 730036 computation_placer.cc:177] computation placer already registered. Please check linkage and avoid linking the same target more than once.
W0000 00:00:1742717801.579284 730036 computation_placer.cc:177] computation placer already registered. Please check linkage and avoid linking the same target more than once.
2025-03-23 13:46:41.581644: I tensorflow/core/platform/cpu_feature_guard.cc:210] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.
To enable the following instructions: AVX2 AVX_VNNI FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.
2025-03-23 13:46:43.120922: E external/local_xla/xla/stream_executor/cuda/cuda_platform.cc:51] failed call to cuInit: INTERNAL: CUDA error: Failed call to cuInit: CUDA_ERROR_UNKNOWN: unknown error
2025-03-23 13:46:43.120943: I external/local_xla/xla/stream_executor/cuda/cuda_diagnostics.cc:178] verbose logging is disabled. Rerun with verbose logging (usually --v=1 or --vmodule=cuda_diagnostics=1) to get more diagnostic output from this module
2025-03-23 13:46:43.120947: I external/local_xla/xla/stream_executor/cuda/cuda_diagnostics.cc:183] retrieving CUDA diagnostic information for host: hitloop
2025-03-23 13:46:43.120948: I external/local_xla/xla/stream_executor/cuda/cuda_diagnostics.cc:190] hostname: hitloop
2025-03-23 13:46:43.121018: I external/local_xla/xla/stream_executor/cuda/cuda_diagnostics.cc:197] libcuda reported version is: 535.230.2
2025-03-23 13:46:43.121024: I external/local_xla/xla/stream_executor/cuda/cuda_diagnostics.cc:201] kernel reported version is: 535.230.2
2025-03-23 13:46:43.121027: I external/local_xla/xla/stream_executor/cuda/cuda_diagnostics.cc:291] kernel version seems to match DS0: 535.230.2
INFO: Started server process [730036]
INFO: Waiting for application startup.
INFO: Application startup complete.
1/1 ██████████ 0s 51ms/step

--- Request Metrics ---
Prediction: 8
Latency (seconds): 0.080555

INFO: 127.0.0.1:58786 - "GET /predict/ HTTP/1.1" 200 OK
1/1 ██████████ 0s 16ms/step

--- Request Metrics ---
Prediction: 8
Latency (seconds): 0.041552


INFO: 127.0.0.1:33658 - "GET /predict/ HTTP/1.1" 200 OK
1/1 ██████████ 0s 16ms/step

--- Request Metrics ---
Prediction: 8
Latency (seconds): 0.032083


INFO: 127.0.0.1:33664 - "GET /predict/ HTTP/1.1" 200 OK
1/1 ██████████ 0s 18ms/step

--- Request Metrics ---
Prediction: 8
Latency (seconds): 0.040463
```

Prediction: 8

1/1  0s 16ms/step

Prediction: 8

1/1  0s 16ms/step

Prediction: 8

1/1 0s 17ms/step

```
--- Request Metrics ---
```

Latency (seconds)

TNEO • 127

```
Prediction: 6
```

Latency (second)

INFO: 127

```

--- Request Metrics ---

```

```
Prediction: 8
```

Latency (second)

INFO: 127

Request Metrics

```
Prediction: 8
```

Efficiency (second)

1/1

```
--- Request Metrics ---
Resolution: 0
```

Latency (seconds): 0.043395

INFO: 127


```
Activities Terminal Mar 23 16:15 samara@hitloop: ~/Videos
INFO: 127.0.0.1:48956 - "GET /predict/ HTTP/1.1" 200 OK
1/1 0s 16ms/step
--- Request Metrics ---
Prediction: 8
Latency (seconds): 0.032006
INFO: 127.0.0.1:48966 - "GET /predict/ HTTP/1.1" 200 OK
1/1 0s 18ms/step
--- Request Metrics ---
Prediction: 8
Latency (seconds): 0.035805
INFO: 127.0.0.1:48974 - "GET /predict/ HTTP/1.1" 200 OK
1/1 0s 18ms/step
--- Request Metrics ---
Prediction: 8
Latency (seconds): 0.043395
INFO: 127.0.0.1:48984 - "GET /predict/ HTTP/1.1" 200 OK
WARNING: WatchFiles detected changes in 'M4_MLOps_Assignment2_Group92.py'. Reloading...
INFO: Shutting down
INFO: Waiting for application shutdown.
INFO: Application shutdown complete.
INFO: Finished server process [730036]
2025-03-23 13:48:40.050839: I tensorflow/core/util/port.cc:153] oneDNN custom operations are on. You may see slightly different numerical results due to floating-point round-off errors from different computation orders. To turn them off, set the environment variable `TF_ENABLE_ONEDNN_OPTS=0`.
2025-03-23 13:48:40.058824: E external/local_xla/xla/stream_executor/cuda/cuda_fft.cc:467] Unable to register cuFFT factory: Attempting to register factory for plugin cuFFT when one has already been registered
WARNING: All log messages before absl::InitializeLog() is called are written to STDERR
E0000 00:00:1742717920.067823 731730 cuda_dnn.cc:8579] Unable to register cuDNN factory: Attempting to register factory for plugin cuDNN when one has already been registered
E0000 00:00:1742717920.070893 731730 cuda_blas.cc:1407] Unable to register cuBLAS factory: Attempting to register factory for plugin cuBLAS when one has already been registered
W0000 00:00:1742717920.078380 731730 computation_placer.cc:177] computation placer already registered. Please check linkage and avoid linking the same target more than once.
W0000 00:00:1742717920.078390 731730 computation_placer.cc:177] computation placer already registered. Please check linkage and avoid linking the same target more than once.
W0000 00:00:1742717920.078391 731730 computation_placer.cc:177] computation placer already registered. Please check linkage and avoid linking the same target more than once.
W0000 00:00:1742717920.078392 731730 computation_placer.cc:177] computation placer already registered. Please check linkage and avoid linking the same target more than once.
2025-03-23 13:48:40.080555: I tensorflow/core/platform/cpu_feature_guard.cc:210] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.
To enable the following instructions: AVX2 AVX_VNNI FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.
2025-03-23 13:48:43.204194: E external/local_xla/xla/stream_executor/cuda/cuda_platform.cc:51] failed call to cuInit: INTERNAL: CUDA error: Failed call to cuInit: CUDA_ERROR_UNKNOWN: unknown error
2025-03-23 13:48:43.204210: I external/local_xla/xla/stream_executor/cuda/cuda_diagnostics.cc:178] verbose logging is disabled. Rerun with verbose logging (usually --v=1 or --vmodule=cuda_diagnostics=1) to get more diagnostic output from this module
2025-03-23 13:48:43.204215: I external/local_xla/xla/stream_executor/cuda/cuda_diagnostics.cc:183] retrieving CUDA diagnostic information for host: hitloop
2025-03-23 13:48:43.204218: I external/local_xla/xla/stream_executor/cuda/cuda_diagnostics.cc:190] hostname: hitloop
2025-03-23 13:48:43.204397: I external/local_xla/xla/stream_executor/cuda/cuda_diagnostics.cc:197] libcuda reported version is: 535.230.2
2025-03-23 13:48:43.204404: I external/local_xla/xla/stream_executor/cuda/cuda_diagnostics.cc:201] kernel reported version is: 535.230.2
2025-03-23 13:48:43.204405: I external/local_xla/xla/stream_executor/cuda/cuda_diagnostics.cc:291] kernel version seems to match DS0: 535.230.2
INFO: Started server process [731730]
INFO: Waiting for application startup.
INFO: Application startup complete.
INFO: 127.0.0.1:38042 - "GET /metrics HTTP/1.1" 200 OK
INFO: 127.0.0.1:51932 - "GET /metrics HTTP/1.1" 200 OK
INFO: 127.0.0.1:38052 - "GET /metrics HTTP/1.1" 200 OK
```

```
(base) samara@hitloop:~/Videos$ for i in {1..10}; do
    echo "Request $i:"
    curl -s http://127.0.0.1:8000/predict/
    echo -e "\n"
done
Request 1:
{"prediction":8,"latency_seconds":0.041551828384399414}

Request 2:
{"prediction":8,"latency_seconds":0.03208279609680176}

Request 3:
{"prediction":8,"latency_seconds":0.04046344757080078}

Request 4:
{"prediction":8,"latency_seconds":0.03537321090698242}

Request 5:
{"prediction":8,"latency_seconds":0.035559892654418945}

Request 6:
{"prediction":8,"latency_seconds":0.03226470947265625}

Request 7:
{"prediction":6,"latency_seconds":0.03362774848937988}

Request 8:
{"prediction":8,"latency_seconds":0.03200554847717285}

Request 9:
{"prediction":8,"latency_seconds":0.03580498695373535}

Request 10:
{"prediction":8,"latency_seconds":0.04339456558227539}

(base) samara@hitloop:~/Videos$
```


Activities Google Chrome Mar 23 16:15

Inbox - 2023 x Home - Goog x M2_Justifica x M3_MLOps_ x M5_MLOps_ x Assignment x Fashion_mnis x Fashion MNIS x lime_explana x MLflow x 127.0.0.1:8000 x + -

127.0.0.1:8000/metrics

CA extra aptitude @sites atten SB CV fco DeepSeek - Int... hl ss v ga PL CSE Python job pr BITS AMS IEEE Xplore Di... moodle

```
# HELP python_gc_objects_collected_total Objects collected during gc
# TYPE python_gc_objects_collected_total counter
python_gc_objects_collected_total{generation="0"} 2942.0
python_gc_objects_collected_total{generation="1"} 700.0
python_gc_objects_collected_total{generation="2"} 10.0
# HELP python_gc_objects_uncollectable_total Uncollectable objects found during GC
# TYPE python_gc_objects_uncollectable_total counter
python_gc_objects_uncollectable_total{generation="0"} 0.0
python_gc_objects_uncollectable_total{generation="1"} 0.0
python_gc_objects_uncollectable_total{generation="2"} 0.0
# HELP python_gc_collections_total Number of times this generation was collected
# TYPE python_gc_collections_total counter
python_gc_collections_total{generation="0"} 1144.0
python_gc_collections_total{generation="1"} 104.0
python_gc_collections_total{generation="2"} 7.0
# HELP python_info Python platform information
# TYPE python_info gauge
python_info{implementation="CPython",major="3",minor="11",patchlevel="11",version="3.11.11"} 1.0
# HELP process_virtual_memory_bytes Virtual memory size in bytes.
# TYPE process_virtual_memory_bytes gauge
process_virtual_memory_bytes 1.0106650624e+010
# HELP process_resident_memory_bytes Resident memory size in bytes.
# TYPE process_resident_memory_bytes gauge
process_resident_memory_bytes 8.55031808e+08
# HELP process_start_time_seconds Start time of the process since unix epoch in seconds.
# TYPE process_start_time_seconds gauge
process_start_time_seconds 1.74271791954e+09
# HELP process_cpu_seconds_total Total user and system CPU time spent in seconds.
# TYPE process_cpu_seconds_total counter
process_cpu_seconds_total 0.91
# HELP process_open_fds Number of open file descriptors.
# TYPE process_open_fds gauge
process_open_fds 23.0
# HELP process_max_fds Maximum number of open file descriptors.
# TYPE process_max_fds gauge
process_max_fds 1024.0
# HELP http_requests_total Total number of requests by method, status and handler.
# TYPE http_requests_total counter
http_requests_total{handler="/metrics",method="GET",status="2xx"} 6.0
# HELP http_requests_created Total number of requests by method, status and handler.
# TYPE http_requests_created gauge
http_requests_created{handler="/metrics",method="GET",status="2xx"} 1.7427180212031662e+09
# HELP http_request_size_bytes Content length of incoming requests by handler. Only value of header is respected. Otherwise ignored. No percentile calculated.
# TYPE http_request_size_bytes summary
http_request_size_bytes_count{handler="/metrics"} 6.0
http_request_size_bytes_sum{handler="/metrics"} 0.0
# HELP http_request_size_bytes_created Content length of incoming requests by handler. Only value of header is respected. Otherwise ignored. No percentile calculated.
# TYPE http_request_size_bytes_created gauge
http_request_size_bytes_created{handler="/metrics"} 1.7427180212032025e+09
# HELP http_response_size_bytes Content length of outgoing responses by handler. Only value of header is respected. Otherwise ignored. No percentile calculated.
# TYPE http_response_size_bytes summary
http_response_size_bytes_count{handler="/metrics"} 6.0
http_response_size_bytes_sum{handler="/metrics"} 36123.0
# HELP http_response_size_bytes_created Content length of outgoing responses by handler. Only value of header is respected. Otherwise ignored. No percentile calculated.
# TYPE http_response_size_bytes_created gauge
http_response_size_bytes_created{handler="/metrics"} 1.7427180212032387e+09
# HELP http_request_duration_highr_seconds Latency with many buckets but no API specific labels. Made for more accurate percentile calculations.
# TYPE http_request_duration_highr_seconds histogram
http_request_duration_highr_seconds_bucket{le="0.01"} 6.0
http_request_duration_highr_seconds_bucket{le="0.025"} 6.0
http_request_duration_highr_seconds_bucket{le="0.05"} 6.0
http_request_duration_highr_seconds_bucket{le="0.075"} 6.0
http_request_duration_highr_seconds_bucket{le="0.1"} 6.0
http_request_duration_highr_seconds_bucket{le="0.25"} 6.0
http_request_duration_highr_seconds_bucket{le="0.5"} 6.0
http_request_duration_highr_seconds_bucket{le="0.75"} 6.0
http_request_duration_highr_seconds_bucket{le="1.0"} 6.0
http_request_duration_highr_seconds_bucket{le="1.5"} 6.0
http_request_duration_highr_seconds_bucket{le="2.0"} 6.0
http_request_duration_highr_seconds_bucket{le="2.5"} 6.0
http_request_duration_highr_seconds_bucket{le="3.0"} 6.0
http_request_duration_highr_seconds_bucket{le="3.5"} 6.0
http_request_duration_highr_seconds_bucket{le="4.0"} 6.0
http_request_duration_highr_seconds_bucket{le="4.5"} 6.0
```

Activities Google Chrome Mar 23 16:15

Inbox - 2023 x Home - Goog x M2_Justifica x M3_MLOps_ x M5_MLOps_ x Assignment x Fashion_mnis x Fashion MNIS x lime_explana x MLflow x 127.0.0.1:8000 x + -

127.0.0.1:8000/metrics

CA extra aptitude @sites atten SB CV fco DeepSeek-Int... hl ss v ga PL CSE Python job pr BITS AMS IEEE Xplore Di... moodle

```
# TYPE process_resident_memory_bytes gauge
process_resident_memory_bytes 8.55031808e+08
# HELP process_start_time_seconds Start time of the process since unix epoch in seconds.
# TYPE process_start_time_seconds gauge
process_start_time_seconds 1.74271791954e+09
# HELP process_cpu_seconds_total Total user and system CPU time spent in seconds.
# TYPE process_cpu_seconds_total counter
process_cpu_seconds_total 9.91
# HELP process_open_fds Number of open file descriptors.
# TYPE process_open_fds gauge
process_open_fds 23.0
# HELP process_max_fds Maximum number of open file descriptors.
# TYPE process_max_fds gauge
process_max_fds 1024.0
# HELP http_requests_total Total number of requests by method, status and handler.
# TYPE http_requests_total counter
http_requests_total{handler="/metrics",method="GET",status="2xx"} 6.0
# HELP http_requests_created Total number of requests by method, status and handler.
# TYPE http_requests_created gauge
http_requests_created{handler="/metrics",method="GET",status="2xx"} 1.7427180212031662e+09
# HELP http_request_size_bytes Content length of incoming requests by handler. Only value of header is respected. Otherwise ignored. No percentile calculated.
# TYPE http_request_size_bytes summary
http_request_size_bytes_count{handler="/metrics"} 6.0
http_request_size_bytes_sum{handler="/metrics"} 0.0
# HELP http_request_size_bytes_created Content length of incoming requests by handler. Only value of header is respected. Otherwise ignored. No percentile calculated.
# TYPE http_request_size_bytes_created gauge
http_request_size_bytes_created{handler="/metrics"} 1.7427180212032025e+09
# HELP http_response_size_bytes Content length of outgoing responses by handler. Only value of header is respected. Otherwise ignored. No percentile calculated.
# TYPE http_response_size_bytes summary
http_response_size_bytes_count{handler="/metrics"} 6.0
http_response_size_bytes_sum{handler="/metrics"} 36123.0
# HELP http_response_size_bytes_created Content length of outgoing responses by handler. Only value of header is respected. Otherwise ignored. No percentile calculated.
# TYPE http_response_size_bytes_created gauge
http_response_size_bytes_created{handler="/metrics"} 1.7427180212032387e+09
# HELP http_request_duration_highr_seconds Latency with many buckets but no API specific labels. Made for more accurate percentile calculations.
# TYPE http_request_duration_highr_seconds histogram
http_request_duration_highr_seconds_bucket{le="0.01"} 6.0
http_request_duration_highr_seconds_bucket{le="0.025"} 6.0
http_request_duration_highr_seconds_bucket{le="0.05"} 6.0
http_request_duration_highr_seconds_bucket{le="0.075"} 6.0
http_request_duration_highr_seconds_bucket{le="0.1"} 6.0
http_request_duration_highr_seconds_bucket{le="0.25"} 6.0
http_request_duration_highr_seconds_bucket{le="0.5"} 6.0
http_request_duration_highr_seconds_bucket{le="0.75"} 6.0
http_request_duration_highr_seconds_bucket{le="1.0"} 6.0
http_request_duration_highr_seconds_bucket{le="1.5"} 6.0
http_request_duration_highr_seconds_bucket{le="2.0"} 6.0
http_request_duration_highr_seconds_bucket{le="2.5"} 6.0
http_request_duration_highr_seconds_bucket{le="3.0"} 6.0
http_request_duration_highr_seconds_bucket{le="3.5"} 6.0
http_request_duration_highr_seconds_bucket{le="4.0"} 6.0
http_request_duration_highr_seconds_bucket{le="4.5"} 6.0
http_request_duration_highr_seconds_bucket{le="5.0"} 6.0
http_request_duration_highr_seconds_bucket{le="7.5"} 6.0
http_request_duration_highr_seconds_bucket{le="10.0"} 6.0
http_request_duration_highr_seconds_bucket{le="30.0"} 6.0
http_request_duration_highr_seconds_bucket{le="60.0"} 6.0
http_request_duration_highr_seconds_bucket{le="+Inf"} 6.0
http_request_duration_highr_seconds_count 6.0
http_request_duration_highr_seconds_sum 0.01723452100122813
# HELP http_request_duration_highr_seconds_created Latency with many buckets but no API specific labels. Made for more accurate percentile calculations.
# TYPE http_request_duration_highr_seconds_created gauge
http_request_duration_highr_seconds_created 1.7427179232813005e+09
# HELP http_request_duration_seconds Latency with only few buckets by handler. Made to be only used if aggregation by handler is important.
# TYPE http_request_duration_seconds histogram
http_request_duration_seconds_bucket{handler="/metrics",le="0.1",method="GET"} 6.0
http_request_duration_seconds_bucket{handler="/metrics",le="0.5",method="GET"} 6.0
http_request_duration_seconds_bucket{handler="/metrics",le="1.0",method="GET"} 6.0
http_request_duration_seconds_bucket{handler="/metrics",le="+Inf",method="GET"} 6.0
http_request_duration_seconds_count{handler="/metrics",method="GET"} 6.0
http_request_duration_seconds_sum{handler="/metrics",method="GET"} 0.01723452100122813
# HELP http_request_duration_seconds_created Latency with only few buckets by handler. Made to be only used if aggregation by handler is important.
# TYPE http_request_duration_seconds_created gauge
http_request_duration_seconds_created{handler="/metrics",method="GET"} 1.7427180212032714e+09
```



```
Activities Terminal Mar 23 16:09 samara@hitloop: ~/Videos
samara@hitloop: ~/Videos samara@hitloop: ~/Videos samara@hitloop: ~/Videos samara@hitloop: ~/Videos
(base) samara@hitloop: ~/Videos$ python M4_MLOps_Assignment2_Group92.py
2025-03-23 13:50:05.649667: I tensorflow/core/util/port.cc:153] oneDNN custom operations are on. You may see slightly different numerical results due to floating-point round-off errors from different computation orders. To turn them off, set the environment variable 'TF_ENABLE_ONEDNN_OPTS=0'.
2025-03-23 13:50:05.657982: E external/local_xla/xla/stream_executor/cuda/cuda_fft.cc:467] Unable to register cuFFT factory: Attempting to register factory for plugin cuFFT when one has already been registered
WARNING: All log messages before absl::InitializeLog() is called are written to STDERR
E0000 00:00:1742718005.668057 732287 cuda_dnn.cc:8579] Unable to register cuDNN factory: Attempting to register factory for plugin cuDNN when one has already been registered
E0000 00:00:1742718005.671385 732287 cuda_blas.cc:1407] Unable to register cuBLAS factory: Attempting to register factory for plugin cuBLAS when one has already been registered
W0000 00:00:1742718005.679872 732287 computation_placer.cc:177] computation placer already registered. Please check linkage and avoid linking the same target more than once.
W0000 00:00:1742718005.679883 732287 computation_placer.cc:177] computation placer already registered. Please check linkage and avoid linking the same target more than once.
W0000 00:00:1742718005.679885 732287 computation_placer.cc:177] computation placer already registered. Please check linkage and avoid linking the same target more than once.
W0000 00:00:1742718005.679887 732287 computation_placer.cc:177] computation placer already registered. Please check linkage and avoid linking the same target more than once.
2025-03-23 13:50:05.682289: I tensorflow/core/platform/cpu_feature_guard.cc:210] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.
To enable the following instructions: AVX2 AVX_VNNI FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.
Starting initial model training...

Training Model: model_run_1 | Epochs: 10 | Batch Size: 64
/home/samara/anaconda3/lib/python3.11/site-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.
  super().__init__(activity_regularizer=activity_regularizer, **kwargs)
2025-03-23 13:50:08.809707: E external/local_xla/xla/stream_executor/cuda/cuda_platform.cc:51] failed call to cuInit: INTERNAL: CUDA error: Failed call to cuInit: CUDA_ERROR_UNKNOWN: unknown error
2025-03-23 13:50:08.809731: I external/local_xla/xla/stream_executor/cuda/cuda_diagnostics.cc:178] verbose logging is disabled. Rerun with verbose logging (usually --v=1 or --vmodule=cuda_diagnostics=1) to get more diagnostic output from this module
2025-03-23 13:50:08.809734: I external/local_xla/xla/stream_executor/cuda/cuda_diagnostics.cc:183] retrieving CUDA diagnostic information for host: hitloop
2025-03-23 13:50:08.809737: I external/local_xla/xla/stream_executor/cuda/cuda_diagnostics.cc:190] hostname: hitloop
2025-03-23 13:50:08.809874: I external/local_xla/xla/stream_executor/cuda/cuda_diagnostics.cc:197] libcuda reported version is: 535.230.2
2025-03-23 13:50:08.809882: I external/local_xla/xla/stream_executor/cuda/cuda_diagnostics.cc:201] kernel reported version is: 535.230.2
2025-03-23 13:50:08.809884: I external/local_xla/xla/stream_executor/cuda/cuda_diagnostics.cc:291] kernel version seems to match DS0: 535.230.2
Epoch 1/10
938/938 — 5s 5ms/step - accuracy: 0.8032 - loss: 0.5680 - val_accuracy: 0.8865 - val_loss: 0.3200
Epoch 2/10
938/938 — 5s 5ms/step - accuracy: 0.8971 - loss: 0.2874 - val_accuracy: 0.8981 - val_loss: 0.2797
Epoch 3/10
938/938 — 5s 5ms/step - accuracy: 0.9117 - loss: 0.2453 - val_accuracy: 0.9032 - val_loss: 0.2645
Epoch 4/10
938/938 — 5s 5ms/step - accuracy: 0.9245 - loss: 0.2080 - val_accuracy: 0.9106 - val_loss: 0.2516
Epoch 5/10
938/938 — 4s 5ms/step - accuracy: 0.9326 - loss: 0.1808 - val_accuracy: 0.9102 - val_loss: 0.2471
Epoch 6/10
938/938 — 4s 5ms/step - accuracy: 0.9432 - loss: 0.1596 - val_accuracy: 0.9131 - val_loss: 0.2573
Epoch 7/10
938/938 — 5s 5ms/step - accuracy: 0.9497 - loss: 0.1386 - val_accuracy: 0.9097 - val_loss: 0.2704
Epoch 8/10
938/938 — 4s 5ms/step - accuracy: 0.9567 - loss: 0.1199 - val_accuracy: 0.9121 - val_loss: 0.2734
Epoch 9/10
938/938 — 5s 5ms/step - accuracy: 0.9611 - loss: 0.1067 - val_accuracy: 0.9142 - val_loss: 0.2671
Epoch 10/10
938/938 — 4s 5ms/step - accuracy: 0.9657 - loss: 0.0944 - val_accuracy: 0.9166 - val_loss: 0.2754
313/313 — 1s 2ms/step
Test Accuracy for model_run_1: 0.9166
Confusion Matrix:
[[892  0 13 16  3  1 70  1  4  0]
 [ 2 983  0 10  3  0  1  0  1  0]
 [ 0  0 999  0  0  0  0  0  0  0]
 [ 0  0  0 999  0  0  0  0  0  0]
 [ 0  0  0  0 999  0  0  0  0  0]
 [ 0  0  0  0  0 999  0  0  0  0]
 [ 0  0  0  0  0  0 999  0  0  0]
 [ 0  0  0  0  0  0  0 999  0  0]
 [ 0  0  0  0  0  0  0  0 999  0]
 [ 0  0  0  0  0  0  0  0  0 999]]
```



```
Activities Terminal Mar 23 16:09 samara@hitloop: ~/Videos
samara@hitloop: ~/Videos samara@hitloop: ~/Videos samara@hitloop: ~/Videos samara@hitloop: ~/Videos
Epoch 16/20
1875/1875 8s 4ms/step - accuracy: 0.9901 - loss: 0.0271 - val_accuracy: 0.9157 - val_loss: 0.4489
Epoch 17/20
1875/1875 8s 4ms/step - accuracy: 0.9926 - loss: 0.0226 - val_accuracy: 0.9151 - val_loss: 0.4590
Epoch 18/20
1875/1875 8s 4ms/step - accuracy: 0.9949 - loss: 0.0173 - val_accuracy: 0.9147 - val_loss: 0.4971
Epoch 19/20
1875/1875 7s 4ms/step - accuracy: 0.9935 - loss: 0.0197 - val_accuracy: 0.9131 - val_loss: 0.5485
Epoch 20/20
1875/1875 8s 4ms/step - accuracy: 0.9941 - loss: 0.0173 - val_accuracy: 0.9136 - val_loss: 0.5196
313/313 1s 2ms/step
Test Accuracy for model_run_2: 0.9136
Confusion Matrix:
[[881  0 18 16  5  1 68  0 11  0]
 [  4 981  0  9  5  0  0  0  1  0]
 [ 23  1 892  6 41  0 35  0  2  0]
 [ 18  2 12 916 25  0 22  0  5  0]
 [  6  1 66 22 858  0 46  0  1  0]
 [  1  0  0  0  0 972  0 16  1 10]
 [116  2 70 18 64  0 719  0 10  1]
 [  0  0  0  0  0  7  0 966  1 26]
 [  3  0  3  6  0  3  2  1 982  0]
 [  1  0  0  0  0  7  0 23  0 969]]

Training Model: model_run_3 | Epochs: 30 | Batch Size: 128
/home/samara/anaconda3/lib/python3.11/site-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.
  super().__init__(activity_regularizer=activity_regularizer, **kwargs)
Epoch 1/30
469/469 3s 6ms/step - accuracy: 0.7839 - loss: 0.6468 - val_accuracy: 0.8750 - val_loss: 0.3560
Epoch 2/30
469/469 3s 6ms/step - accuracy: 0.8895 - loss: 0.3110 - val_accuracy: 0.8786 - val_loss: 0.3257
Epoch 3/30
469/469 3s 6ms/step - accuracy: 0.9067 - loss: 0.2631 - val_accuracy: 0.8984 - val_loss: 0.2757
Epoch 4/30
469/469 3s 7ms/step - accuracy: 0.9151 - loss: 0.2341 - val_accuracy: 0.9033 - val_loss: 0.2642
Epoch 5/30
469/469 3s 6ms/step - accuracy: 0.9251 - loss: 0.2076 - val_accuracy: 0.9043 - val_loss: 0.2674
Epoch 6/30
469/469 3s 7ms/step - accuracy: 0.9322 - loss: 0.1841 - val_accuracy: 0.9106 - val_loss: 0.2521
Epoch 7/30
469/469 3s 6ms/step - accuracy: 0.9392 - loss: 0.1650 - val_accuracy: 0.9102 - val_loss: 0.2566
Epoch 8/30
469/469 3s 6ms/step - accuracy: 0.9475 - loss: 0.1470 - val_accuracy: 0.9137 - val_loss: 0.2445
Epoch 9/30
469/469 3s 7ms/step - accuracy: 0.9508 - loss: 0.1365 - val_accuracy: 0.9150 - val_loss: 0.2451
Epoch 10/30
469/469 3s 7ms/step - accuracy: 0.9577 - loss: 0.1193 - val_accuracy: 0.9118 - val_loss: 0.2652
Epoch 11/30
469/469 3s 7ms/step - accuracy: 0.9617 - loss: 0.1092 - val_accuracy: 0.9172 - val_loss: 0.2542
Epoch 12/30
469/469 3s 7ms/step - accuracy: 0.9634 - loss: 0.0995 - val_accuracy: 0.9199 - val_loss: 0.2669
Epoch 13/30
469/469 3s 6ms/step - accuracy: 0.9686 - loss: 0.0896 - val_accuracy: 0.9165 - val_loss: 0.2684
Epoch 14/30
```

ActivitiesTerminalMar 23 16:09samara@hitloop: ~/Videos

Epoch 26/30
469/469 3s 7ms/step - accuracy: 0.9949 - loss: 0.0173 - val_accuracy: 0.9133 - val_loss: 0.4470
Epoch 27/30
469/469 3s 7ms/step - accuracy: 0.9935 - loss: 0.0199 - val_accuracy: 0.9154 - val_loss: 0.4498
Epoch 28/30
469/469 3s 7ms/step - accuracy: 0.9935 - loss: 0.0202 - val_accuracy: 0.9109 - val_loss: 0.4581
Epoch 29/30
469/469 3s 7ms/step - accuracy: 0.9941 - loss: 0.0181 - val_accuracy: 0.9162 - val_loss: 0.4549
Epoch 30/30
469/469 3s 7ms/step - accuracy: 0.9963 - loss: 0.0127 - val_accuracy: 0.9151 - val_loss: 0.4975
313/313 1s 2ms/step
Test Accuracy for model_run_3: 0.9151
Confusion Matrix:
[[915 0 18 13 4 2 42 0 6 0]
[2 983 1 13 0 0 0 0 1 0]
[21 1 875 9 55 0 38 0 0 1]
[22 4 7 951 7 0 9 0 0 0]
[5 1 52 45 869 0 26 0 2 0]
[2 0 0 1 0 977 0 13 0 7]
[169 0 62 29 69 0 666 0 5 0]
[0 0 0 0 0 6 0 976 1 17]
[5 1 5 6 0 2 5 2 973 1]
[0 0 0 0 0 6 1 27 0 966]]

Training Model: model_run_4 | Epochs: 20 | Batch Size: 64
/home/samara/anaconda3/lib/python3.11/site-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
Epoch 1/20
938/938 5s 5ms/step - accuracy: 0.8017 - loss: 0.5735 - val_accuracy: 0.8861 - val_loss: 0.3264
Epoch 2/20
938/938 5s 5ms/step - accuracy: 0.9002 - loss: 0.2805 - val_accuracy: 0.8966 - val_loss: 0.2815
Epoch 3/20
938/938 5s 5ms/step - accuracy: 0.9137 - loss: 0.2374 - val_accuracy: 0.9023 - val_loss: 0.2684
Epoch 4/20
938/938 5s 5ms/step - accuracy: 0.9226 - loss: 0.2100 - val_accuracy: 0.9088 - val_loss: 0.2547
Epoch 5/20
938/938 5s 5ms/step - accuracy: 0.9332 - loss: 0.1802 - val_accuracy: 0.9096 - val_loss: 0.2541
Epoch 6/20
938/938 5s 5ms/step - accuracy: 0.9408 - loss: 0.1615 - val_accuracy: 0.9149 - val_loss: 0.2519
Epoch 7/20
938/938 5s 5ms/step - accuracy: 0.9475 - loss: 0.1437 - val_accuracy: 0.9168 - val_loss: 0.2540
Epoch 8/20
938/938 5s 5ms/step - accuracy: 0.9530 - loss: 0.1264 - val_accuracy: 0.9162 - val_loss: 0.2548
Epoch 9/20
938/938 5s 5ms/step - accuracy: 0.9614 - loss: 0.1057 - val_accuracy: 0.9201 - val_loss: 0.2511
Epoch 10/20
938/938 5s 5ms/step - accuracy: 0.9665 - loss: 0.0934 - val_accuracy: 0.9182 - val_loss: 0.2643
Epoch 11/20
938/938 4s 5ms/step - accuracy: 0.9734 - loss: 0.0779 - val_accuracy: 0.9197 - val_loss: 0.2806
Epoch 12/20
938/938 5s 5ms/step - accuracy: 0.9736 - loss: 0.0720 - val_accuracy: 0.9195 - val_loss: 0.2971
Epoch 13/20
938/938 5s 5ms/step - accuracy: 0.9788 - loss: 0.0594 - val_accuracy: 0.9196 - val_loss: 0.3034


```
Activities Terminal Mar 23 16:09 samara@hitloop: ~/Videos
samara@hitloop: ~/Videos samara@hitloop: ~/Videos samara@hitloop: ~/Videos samara@hitloop: ~/Videos
Epoch 15/20
938/938 5s 5ms/step - accuracy: 0.9840 - loss: 0.0456 - val_accuracy: 0.9205 - val_loss: 0.3197
Epoch 16/20
938/938 5s 5ms/step - accuracy: 0.9871 - loss: 0.0386 - val_accuracy: 0.9187 - val_loss: 0.3419
Epoch 17/20
938/938 5s 5ms/step - accuracy: 0.9885 - loss: 0.0339 - val_accuracy: 0.9121 - val_loss: 0.3831
Epoch 18/20
938/938 5s 5ms/step - accuracy: 0.9890 - loss: 0.0306 - val_accuracy: 0.9165 - val_loss: 0.3587
Epoch 19/20
938/938 5s 5ms/step - accuracy: 0.9911 - loss: 0.0265 - val_accuracy: 0.9168 - val_loss: 0.4007
Epoch 20/20
938/938 4s 5ms/step - accuracy: 0.9918 - loss: 0.0246 - val_accuracy: 0.9142 - val_loss: 0.4512
313/313 1s 2ms/step
Test Accuracy for model_run_4: 0.9142
Confusion Matrix:
[[864 0 22 13 5 1 92 0 3 0]
 [ 5 979 1 8 5 0 1 0 1 0]
 [ 13 0 888 6 39 0 53 0 1 0]
 [ 34 2 16 900 21 0 26 0 1 0]
 [ 0 1 77 24 845 0 52 0 1 0]
 [ 0 0 0 0 0 984 0 3 1 12]
 [ 93 0 62 16 40 0 785 0 4 0]
 [ 0 0 0 0 0 20 0 942 0 38]
 [ 3 0 5 5 0 2 5 0 980 0]
 [ 0 0 0 0 0 5 1 18 1 975]]

Training Model: model_run_5 | Epochs: 30 | Batch Size: 32
/home/samara/anaconda3/lib/python3.11/site-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.
  super().__init__(activity_regularizer=activity_regularizer, **kwargs)
Epoch 1/30
1875/1875 8s 4ms/step - accuracy: 0.8134 - loss: 0.5249 - val_accuracy: 0.8828 - val_loss: 0.3260
Epoch 2/30
1875/1875 8s 4ms/step - accuracy: 0.8991 - loss: 0.2773 - val_accuracy: 0.8990 - val_loss: 0.2796
Epoch 3/30
1875/1875 8s 4ms/step - accuracy: 0.9183 - loss: 0.2231 - val_accuracy: 0.8973 - val_loss: 0.2729
Epoch 4/30
1875/1875 8s 4ms/step - accuracy: 0.9330 - loss: 0.1867 - val_accuracy: 0.9074 - val_loss: 0.2655
Epoch 5/30
1875/1875 7s 4ms/step - accuracy: 0.9437 - loss: 0.1545 - val_accuracy: 0.9149 - val_loss: 0.2526
Epoch 6/30
1875/1875 8s 4ms/step - accuracy: 0.9524 - loss: 0.1310 - val_accuracy: 0.9129 - val_loss: 0.2574
Epoch 7/30
1875/1875 8s 4ms/step - accuracy: 0.9585 - loss: 0.1124 - val_accuracy: 0.9135 - val_loss: 0.2610
Epoch 8/30
1875/1875 7s 4ms/step - accuracy: 0.9672 - loss: 0.0923 - val_accuracy: 0.9121 - val_loss: 0.2975
Epoch 9/30
1875/1875 8s 4ms/step - accuracy: 0.9719 - loss: 0.0782 - val_accuracy: 0.9158 - val_loss: 0.2988
Epoch 10/30
1875/1875 8s 4ms/step - accuracy: 0.9764 - loss: 0.0670 - val_accuracy: 0.9138 - val_loss: 0.3195
Epoch 11/30
1875/1875 8s 4ms/step - accuracy: 0.9794 - loss: 0.0583 - val_accuracy: 0.9170 - val_loss: 0.3249
Epoch 12/30
1875/1875 8s 4ms/step - accuracy: 0.9822 - loss: 0.0482 - val_accuracy: 0.9136 - val_loss: 0.3838
Epoch 13/30
```

```
Activities Terminal Mar 23 16:09 samara@hitloop: ~/Videos
samara@hitloop: ~/Videos samara@hitloop: ~/Videos samara@hitloop: ~/Videos samara@hitloop: ~/Videos
1875/1875 8s 4ms/step - accuracy: 0.9822 - loss: 0.0482 - val_accuracy: 0.9136 - val_loss: 0.3838
Epoch 13/30
1875/1875 7s 4ms/step - accuracy: 0.9852 - loss: 0.0422 - val_accuracy: 0.9164 - val_loss: 0.3786
Epoch 14/30
1875/1875 8s 4ms/step - accuracy: 0.9872 - loss: 0.0380 - val_accuracy: 0.9141 - val_loss: 0.4194
Epoch 15/30
1875/1875 8s 4ms/step - accuracy: 0.9893 - loss: 0.0314 - val_accuracy: 0.9132 - val_loss: 0.4315
Epoch 16/30
1875/1875 8s 4ms/step - accuracy: 0.9910 - loss: 0.0255 - val_accuracy: 0.9133 - val_loss: 0.4684
Epoch 17/30
1875/1875 8s 4ms/step - accuracy: 0.9915 - loss: 0.0248 - val_accuracy: 0.9130 - val_loss: 0.4696
Epoch 18/30
1875/1875 8s 4ms/step - accuracy: 0.9935 - loss: 0.0201 - val_accuracy: 0.9109 - val_loss: 0.5068
Epoch 19/30
1875/1875 8s 4ms/step - accuracy: 0.9939 - loss: 0.0187 - val_accuracy: 0.9111 - val_loss: 0.5407
Epoch 20/30
1875/1875 8s 4ms/step - accuracy: 0.9918 - loss: 0.0249 - val_accuracy: 0.9163 - val_loss: 0.5325
Epoch 21/30
1875/1875 8s 4ms/step - accuracy: 0.9953 - loss: 0.0143 - val_accuracy: 0.9130 - val_loss: 0.5556
Epoch 22/30
1875/1875 8s 4ms/step - accuracy: 0.9948 - loss: 0.0153 - val_accuracy: 0.9111 - val_loss: 0.5826
Epoch 23/30
1875/1875 8s 4ms/step - accuracy: 0.9955 - loss: 0.0137 - val_accuracy: 0.9121 - val_loss: 0.5744
Epoch 24/30
1875/1875 8s 4ms/step - accuracy: 0.9948 - loss: 0.0171 - val_accuracy: 0.9064 - val_loss: 0.6283
Epoch 25/30
1875/1875 7s 4ms/step - accuracy: 0.9948 - loss: 0.0172 - val_accuracy: 0.9122 - val_loss: 0.6078
Epoch 26/30
1875/1875 7s 4ms/step - accuracy: 0.9965 - loss: 0.0101 - val_accuracy: 0.9163 - val_loss: 0.6716
Epoch 27/30
1875/1875 8s 4ms/step - accuracy: 0.9954 - loss: 0.0140 - val_accuracy: 0.9146 - val_loss: 0.6234
Epoch 28/30
1875/1875 8s 4ms/step - accuracy: 0.9957 - loss: 0.0133 - val_accuracy: 0.9158 - val_loss: 0.6242
Epoch 29/30
1875/1875 8s 5ms/step - accuracy: 0.9975 - loss: 0.0086 - val_accuracy: 0.9102 - val_loss: 0.6431
Epoch 30/30
1875/1875 8s 4ms/step - accuracy: 0.9967 - loss: 0.0101 - val_accuracy: 0.9093 - val_loss: 0.6929
313/313 1s 2ms/step
Test Accuracy for model_run_5: 0.9093
Confusion Matrix:
[[825  0  38  15  4  2 108  1  6  1]
 [  0 981  3  12  1  0  2  0  1  0]
 [ 10  1 878  5  42  0  64  0  0  0]
 [  9  3  25 905  31  0  26  0  1  0]
 [  1  1  77  14 823  0  82  0  2  0]
 [  0  0  0  0  0 979  0  8  1 12]
 [ 79  1  54  20  41  0 795  0  9  1]
 [  0  0  0  0  0  12  0 957  1 30]
 [  2  1  6  2  2  4  2  0 980  1]
 [  1  0  0  0  0  8  0  21  0 970]]

Best model saved as: best_model.pkl with Accuracy: 0.9166
Loading best model and checking for drift with original and synthetic data...
```


Activities

Google Chrome

Mar 23 16:12

Inbox - 2023

Home - Goog

M2_Justific

M3_MLOps

M5_MLOps

Assignment

fashion_mnist

Fashion MNIS

lime_explaine

MLflow

127.0.0.1:80

127.0.0.1:5000/#/experiments/695082901439436626?searchFilter=&orderByKey=attributes.start_time&orderByAsc=false&startTime=ALL&lifecycleFilter=Active&modelVersionFilter=All+Runs&datasetsFilter=W10%3D

CAextraaptitude@sitesattenSB CV fcoDeepSeek - Int...hlssvgaPL CSE Python job prBITSAMSIEEE Xplore Di...moodle

mlflow2.21.0

ExperimentsModelsPrompts

GitHubDocs

Experiments

Search experiments

Default

fashion_mnist_drift_detection_task

Runs

Evaluation

Experimental

Traces

metrics.rmse < 1 and params.model = "tree"

Time created

State: Active

Datasets

Sort: Created

Columns

Group by

New run

Run Name	Created	Dataset	Duration	Source	Models
Retraining_Alter_Drift	2 hours ago	-	15.0s	M4_ML...	-
Drift_Detection_Artificial	2 hours ago	-	4ms	M4_ML...	-
Drift_Detection_Original	2 hours ago	-	3ms	M4_ML...	-
model_run_5	2 hours ago	-	4.0min	M4_ML...	-
model_run_4	2 hours ago	-	1.6min	M4_ML...	-
model_run_3	2 hours ago	-	1.6min	M4_ML...	-
model_run_2	2 hours ago	-	2.6min	M4_ML...	-
model_run_1	2 hours ago	-	48.2s	M4_ML...	-

8 matching runs

Activities

Google Chrome

Mar 23 16:12

Inbox - 2023

Home - Goog

M2_Justifical

M3_MLOps_

M5_MLOps_

Assignment

fashion_mnist

Fashion MNIS

lime_explana

MLflow

127.0.0.1:80

127.0.0.1:5000/#/experiments/695082901439436626/runs/6e5e529399e644198a06388cfe357e95

CAextraaptitude@sitesattenSB CV fcoDeepSeek - Int...hlssv ga PL CSE Python job pr BITSAMSIEEE Xplore Di...moodle

mlflow2.21.0

ExperimentsModelsPrompts

en

GitHubDocs

fashion_mnist_drift_detection_task

model_run_1

OverviewModel metricsSystem metricsTracesArtifacts

Description

No description

Details

Created at	03/23/2025, 01:50:07 PM
Created by	samara
Experiment ID	695082901439436626
Status	Finished
Run ID	6e5e529399e644198a06388cfe357e95
Duration	48.2s
Datasets used	—
Tags	Add tags
Source	M4_MLOps_Assignment2_Group92.py
Logged models	—
Registered models	—
Registered prompts	—

Parameters (2)

Search parameters

Parameter	Value
batch_size	64
epochs	10

Metrics (4)

Search metrics

Metric	Value
recall	0.9166
accuracy	0.9166
f1_score	0.916611250709106
precision	0.9171038433616341

Activities

Google Chrome

Mar 23 16:12

Inbox - 2023

Home - Goog

M2_Justific

M3_MLOps

M5_MLOps

Assignment

fashion_mnist

Fashion MNIS

lime_explaine

MLflow

127.0.0.1:80

127.0.0.1:5000/#/experiments/695082901439436626/runs/6e5e529399e644198a06388cfe357e95/model-metrics

CAextraaptitude@sitesattenSB CV fcoDeepSeek - Int...hlssvgaPL CSE Python job prBITSAMSIEEE Xplore Di...moodle

mlflow2.21.0

ExperimentsModelsPrompts

GitHubDocs

fashion_mnist_drift_detection_task

model_run_1

Overview

Model metrics

System metrics

Traces

Artifacts


Search metric charts

Auto-refresh

Model metrics (4)

Add chart


accuracy



0.92

model_run_1


f1_score



0.92

model_run_1


precision



0.92

model_run_1

recall



0.92

model_run_1

Add section

Activities

Google Chrome

Mar 23 16:12

Inbox - 2023

Home - Goog

M2_Justifical

M3_MLOps_

M5_MLOps_

Assignment

fashion_mnist

Fashion MNIS

lime_explana

MLflow

127.0.0.1:80

127.0.0.1:5000/#/experiments/695082901439436626/runs/9f3756e14a6a4e53a52e5eec30f710b5

CAextraaptitude@sitesattenSB CV fcoDeepSeek - Int...hlssv ga PL CSE Python job pr BITSAMSIEEE Xplore Di...moodle

mlflow2.21.0

ExperimentsModelsPrompts

en

GitHubDocs

fashion_mnist_drift_detection_task

model_run_2

OverviewModel metricsSystem metricsTracesArtifacts

Description

No description

Details

Created at	03/23/2025, 01:50:55 PM
Created by	samara
Experiment ID	695082901439436626
Status	Finished
Run ID	9f3756e14a6a4e53a52e5eec30f710b5
Duration	2.6min
Datasets used	—
Tags	Add tags
Source	M4_MLOps_Assignment2_Group92.py
Logged models	—
Registered models	—
Registered prompts	—

Parameters (2)

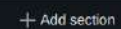
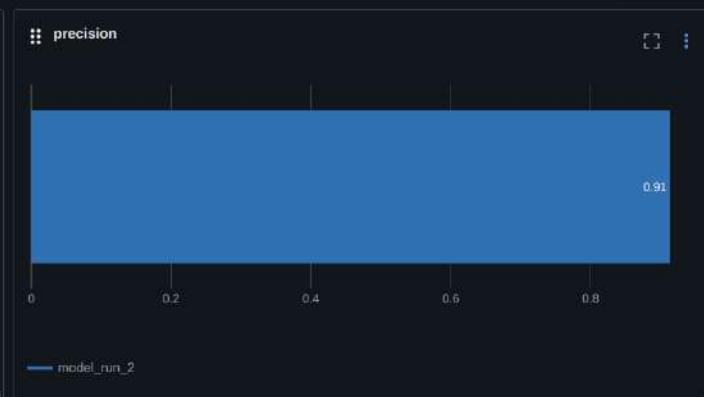
Search parameters

Parameter	Value
batch_size	32
epochs	20

Metrics (4)

Search metrics

Metric	Value
recall	0.9136
accuracy	0.9136
f1_score	0.9130662623626933
precision	0.9133244139893585



Activities

Google Chrome

Mar 23 16:13

Inbox - 2023

Home - Goog

M2_Justific

M3_MLOps/

M5_MLOps/

Assignment

fashion_mnist

Fashion MNIS

lime_explana

MLflow

127.0.0.1:800

+

127.0.0.1:5000/#/experiments/695082901439436626/runs/16ad104e762348879a2a5ac4e246dd0b

CA

extra

aptitude

@sites

atten

SB

CV

fco

DeepSeek - Int...

hl

ss

v

ga

PL

CSE

Python

job

pr

BITS

AMS

IEEE Xplore Di...

moodle

»

mlflow2.21.0

Experiments

Models

Prompts

en

GitHub

Docs

fashion_mnist_drift_detection_task

model_run_3

⋮

Overview

Model metrics

System metrics

Traces

Artifacts

Description

No description

Details

Created at	03/23/2025, 01:53:29 PM
Created by	samara
Experiment ID	695082901439436626
Status	Finished
Run ID	16ad104e762348879a2a5ac4e246dd0b
Duration	1.6min
Datasets used	—
Tags	Add tags
Source	M4_MLOps_Assignment2_Group92.py
Logged models	—
Registered models	—
Registered prompts	—

Parameters (2)

Search parameters

Parameter	Value
batch_size	128
epochs	30

Metrics (4)

Search metrics

Metric	Value
recall	0.9151
accuracy	0.9151
f1_score	0.9138318684657559
precision	0.9157011953506595

Activities

Google Chrome

Mar 23 16:13

Inbox - 2023

Home - Goog

M2_Justifical

M3_MLOps

M5_MLOps

Assignment

fashion_mnist

Fashion MNIST

lime_explaine

MLflow

127.0.0.1:80

127.0.0.1:5000/#/experiments/695082901439436626/runs/16ad104e762348879a2a5ac4e246dd0b/model-metrics

CAextraaptitude@sitesattenSB CV fcoDeepSeek - Int...hlssvgaPL CSE Python job prBITSAMSIEEE Xplore Di...moodle

mlflow2.21.0

ExperimentsModelsPrompts

GitHubDocs

fashion_mnist_drift_detection_task

model_run_3

Overview

Model metrics

System metrics

Traces

Artifacts

Search metric charts


Auto-refresh

Model metrics (4)

+

Add chart


accuracy



0.92

model_run_3


f1_score



0.91

model_run_3

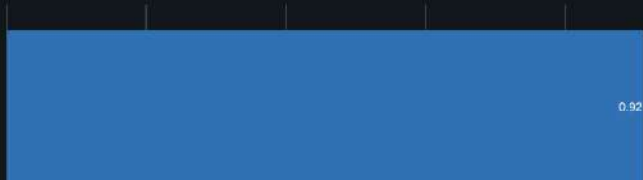
precision



0.92

model_run_3

recall



0.92

model_run_3

+

Add section

Overview **Model metrics** **System metrics** **Traces** **Artifacts**

Description [🔗](#)

Details

Activities

Google Chrome

Mar 23 16:13

Inbox - 2023

Home - Goog

M2_Justific

M3_MLOps

M5_MLOps

Assignment

fashion_mnist

Fashion MNIS

lime_explana

MLflow

127.0.0.1:80

127.0.0.1:5000/#/experiments/695082901439436626/runs/4bf264ad94334a609b4269c6094c49cc/model-metrics

CAextraaptitude@sitesattenSB CV fcoDeepSeek - Int...hlssvgaPL CSE Python job prBITSAMSIEEE Xplore Di...moodle

mlflow2.21.0

ExperimentsModelsPrompts

GitHubDocs

fashion_mnist_drift_detection_task

model_run_4

OverviewModel metricsSystem metricsTracesArtifacts


Search metric charts

Auto-refresh

Model metrics (4)

Add chart

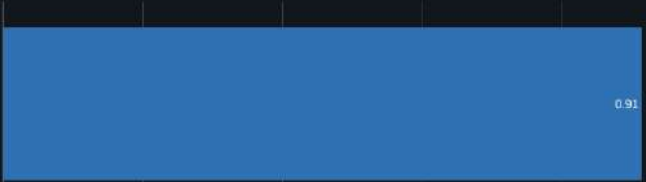
accuracy



0.91

model_run_4


f1_score



0.91

model_run_4


precision



0.92

model_run_4

recall



0.91

model_run_4

Add section

Activities

Google Chrome

Mar 23 16:13

Inbox - 2023

Home - Goog

M2_Justifical

M3_MLOps/

M5_MLOps/

Assignment

fashion_mnist

Fashion MNIS

lime_explana

MLflow

127.0.0.1:80

127.0.0.1:5000/#/experiments/695082901439436626/runs/b982e2dda9774c4c97e1e5cff6a1a651

CAextraaptitude@sitesattenSB CV fcoDeepSeek - Int...hlssvgaPL CSE Python job prBITSAMSIEEE Xplore Di...moodle

mlflow2.21.0

ExperimentsModelsPrompts

enGitHubDocs

fashion_mnist_drift_detection_task

model_run_5

OverviewModel metricsSystem metricsTracesArtifacts

Description

No description

Details

Created at	03/23/2025, 01:56:38 PM
Created by	samara
Experiment ID	695082901439436626
Status	Finished
Run ID	b982e2dda9774c4c97e1e5cff6a1a651
Duration	4.0min
Datasets used	—
Tags	Add tags
Source	M4_MLOps_Assignment2_Group92.py
Logged models	—
Registered models	—
Registered prompts	—

Parameters (2)

Search parameters

Parameter	Value
batch_size	32
epochs	30

Metrics (4)

Search metrics

Metric	Value
recall	0.9093
accuracy	0.9093
f1_score	0.9097992336215888
precision	0.9112242166891751

Activities

Google Chrome

Mar 23 16:13

Inbox - 2023

Home - Goog

M2_Justifical

M3_MLOps

M5_MLOps

Assignment

fashion_mnist

Fashion MNIST

lime_explaine

MLflow

127.0.0.1:80

127.0.0.1:5000/#/experiments/695082901439436626/runs/b982e2dda9774c4c97e1e5cff6a1a651/model-metrics

CA

extra

aptitude

@sites

atten

SB

CV

fco

DeepSeek - Int...

hl

ss

v

ga

PL

CSE

Python

job

pr

BITS

AMS

IEEE Xplore Di...

moodle

mlflow2.21.0

Experiments

Models

Prompts

GitHub

Docs

fashion_mnist_drift_detection_task

model_run_5

Overview

Model metrics

System metrics

Traces

Artifacts

Search metric charts

Auto-refresh

Model metrics (4)

Add chart

accuracy

0.91

model_run_5

f1_score

0.91

model_run_5

precision

0.91

model_run_5

recall

0.91

model_run_5

Add section